

# Introduction to Sockets

# What is a socket?

- The ***socket*** is the BSD method for accomplishing inter-process communication (IPC).
- It is used to allow one process to speak to another (on same or different machine).
  - ***Analogy***: Like the telephone is used to allow one person to speak to another.
- Works very similar to files.
  - Socket descriptor \_very similar to file descriptor.
  - Read/write on a socket and file are very similar.

# Basic Idea

When two processes located on the same or different machines communicate, we define association and socket.

– **Association**: basically a 5-tuple

- Protocol
- Local IP address
- Local port number
- Remote IP address
- Remote port number

– **Socket**: also called half-association (a 3-tuple)

- Protocol, local IP address, local port number
- Protocol, remote IP address, remote port number

# More about sockets

Creating a socket is the first step in network programming using BSD socket interface.

- Using the *socket()* system call.
- Two main addressing formats of a socket:
  - ***AF\_UNIX***: uses Unix pathnames to identify sockets, and are very useful for IPC between processes on the same machine.
  - ***AF\_INET***: uses IP addresses.
- In addition to machine address, there is also a port number that allows more than one *AF\_INET* socket on each machine.

# Types of socket

Two most common types:

- ***SOCK\_STREAM***: Stream sockets, which provide reliable, two-way, connection-oriented communication streams. *<Uses TCP>*
- ***SOCK\_DGRAM***: Datagram sockets, which provide connectionless, unreliable service, used for packet-by-packet transfer of information. *<Uses UDP>*
- Other types like SOCK\_RAW also exist.
  - Beyond the scope of the present discussion.

# Systems calls for using sockets

- `socket()`
- `bind()`
- `connect()`
- `listen()`
- `accept()`
- `send()` & `recv()`
- `sendto()` & `recvfrom()`
- `close()` & `shutdown()`
- `getpeername()`
- `gethostname()`
- `gethostbyname()`

# socket() :: Get the Socket Descriptor

- **General syntax:**

- *domain*: should be set to AF\_INET (typically)
- *type*: should be set to SOCK\_STREAM or SOCK\_DGRAM
- *protocol*: set to zero (typically)

- **Returns**: socket descriptor; -1 on error

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket (int domain, int type, int protocol)
```

# bind():: What Port am I on?

Used to associate the socket with an address

- **General syntax:**

- *sockfd*: socket file descriptor returned by *socket()*
- *my\_addr*: pointer to a structure that contains information about the local IP address and port number.
- *addrlen*: typically set to *sizeof(struct sockaddr)*

- **Returns:** -1 on error

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int bind (int sockfd, struct sockaddr *my_addr, int addrlen);
```



# The *sockaddr* Structure

```
struct sockaddr
```

```
{
```

```
    unsigned short sa_family;
```

```
    char sa_data[14];
```

```
}
```

```
struct sockaddr_in
```

```
{
```

```
    short int sin_family;
```

```
    unsigned short int sin_port;
```

```
    struct in_addr sin_addr;
```

```
    unsigned char sin_zero[8];
```

```
}
```

***sockaddr\_in*** is a parallel structure to ***sockaddr*** which a programmer uses

in the program for convenience.

```
struct in_addr
```

```
{
```

```
    unsigned long s_addr;
```

```
}
```

# connect(): Connect to a Remote Socket

- **General syntax:**

- *sockfd*: socket file descriptor returned by *socket()*
- *serv\_addr*: pointer to a structure that contains the destination IP address and the port number
- *addrlen*: typically set to *sizeof(struct sockaddr)*

- **Returns:** -1 on error

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int connect (int sockfd, struct sockaddr *serv_addr, int addrlen);
```

# listen(): Get Set for Incoming Connections

- Here, we wish to wait for incoming connections and handle them in some way.
    - Two steps, first you *listen()*, then you *accept()*.
  - **General syntax:**
    - *sockfd*: socket file descriptor returned by *socket()*.
    - *backlog*: used to set the maximum number of requests (up to a maximum of about 20) that will be queued up before requests start being denied.
  - **Returns:** -1 on error
- int listen (int sockfd, int backlog);*

# accept(): Waiting for Incoming Connections

- Basic concept:
  - Someone far away will try to *connect()* to your machine on a port that you are *listen()*'ing on.
  - Such connections will be queued up waiting to be *accept()*'ed.
  - *accept()* returns a ***brand new socket file descriptor*** to use for every single connection.
- Two socket file descriptors!!
  - The original one is still listening on your port.
  - Newly created one is finally ready to *send()* and *recv()*.

# accept(): contd..

## ***General syntax:***

- *sockfd*: *listen()*'ing socket descriptor
- *addr*: pointer to a local *struct sockaddr\_in* (This is where the information about the incoming connection will go)
- *addrlen*: local integer variable that should be set to *sizeof(struct sockaddr\_in)* before *accept()* is called.

- ***Returns***: -1 on error

*#include <sys/socket.h>*

*int accept (int sockfd, void \*addr, int \*addrlen);*

# send() and recv(): Sending/receiving Data

- Used for communicating over stream sockets or connected datagram sockets.

- ***General syntax:***

- *mesg*: a pointer to the data you want to send
- *len*: length of the data in bytes
- *buf*: buffer to read the information into
- *flags*: typically set to 0

- *send()* returns the number of bytes actually sent out, and *recv()* returns the number of bytes actually read into the buffer.

*int send (int sockfd, const void \*mesg, int len, int flags);*

*int recv (int sockfd, void \*buf, int len, unsigned int flags);*

# sendto() and recvfrom()

Used to transmit and receive data packets over unconnected datagram sockets.

- ***General syntax:***

If you *connect()* a datagram socket, you can then simply use *send()* and *recv()* for all your transactions.

```
int sendto (int sockfd, const void *msg, int len, unsigned int flags,  
const struct sockaddr *to, int tolen);
```

```
int recvfrom (int sockfd, void *buf, int len, unsigned int flags,  
struct sockaddr *from, int *fromlen);
```

# close() and shutdown()

Used to close the connection on the socket descriptor.

- This prevents any more reads and writes to the socket.
- `how=0` \_further receives are disallowed
- `how=1` \_further sends are disallowed
- `how=2` \_further sends and receives are disallowed (like *close()*)

*close (sockfd);*

*int shutdown (int sockfd, int how);*



# getpeername()

- This function will tell you who is at the other end of a connection stream socket.
  - *sockfd*: descriptor of the connected stream socket
  - *addr*: pointer to a structure that will hold the information about the other side of the connection
  - *addrlen*: pointer to an *int* that should be initialized to *sizeof(struct sockaddr)*

```
#include <sys/socket.h>
```

```
int getpeername (int sockfd, struct sockaddr *addr, int *addrlen);
```

# gethostname()

This function returns the name of the computer that your program is running on.

- This name can be used by *gethostbyname()* to determine the IP address of the local machine.
- *hostname*: pointer to an array of *chars* that will contain the host name upon the function's return.
- *size*: length in bytes of the *hostname* array.

```
#include <unistd.h>
```

```
int gethostname (char *hostname, size_t size);
```

# gethostbyname()

Returns the IP address of a host given its name.

- Invokes the *Domain Name Server (DNS)*.
- Returns a pointer to a struct hostent:

```
#include <netdb.h>

struct hostent *gethostbyname (const char *name);

struct hostent
{
    char *h_name; /* official name of the host */
    char **h_aliases; /* NULL terminate array of alternate names */
    int h_addrtype; /* Type of address being returned (AF_INET) */
    int h_length; /* Length of the address in bytes */
    char **h_addr_list; /* Zero terminated array of network addresses */
};

#define h_addr h_addr_list[0];
```

# Client-server Model

Standard model for network applications.

– A **server** is a process that is waiting to be contacted by a **client** process so as to provide some service.

- **Typical scenario:**

- The server process is started on some computer system.

- Initializes itself, then goes to sleep waiting for a client request.

- A client process is started, either on the same system or on some other system.

- Client sends a request (across the network) to the server.

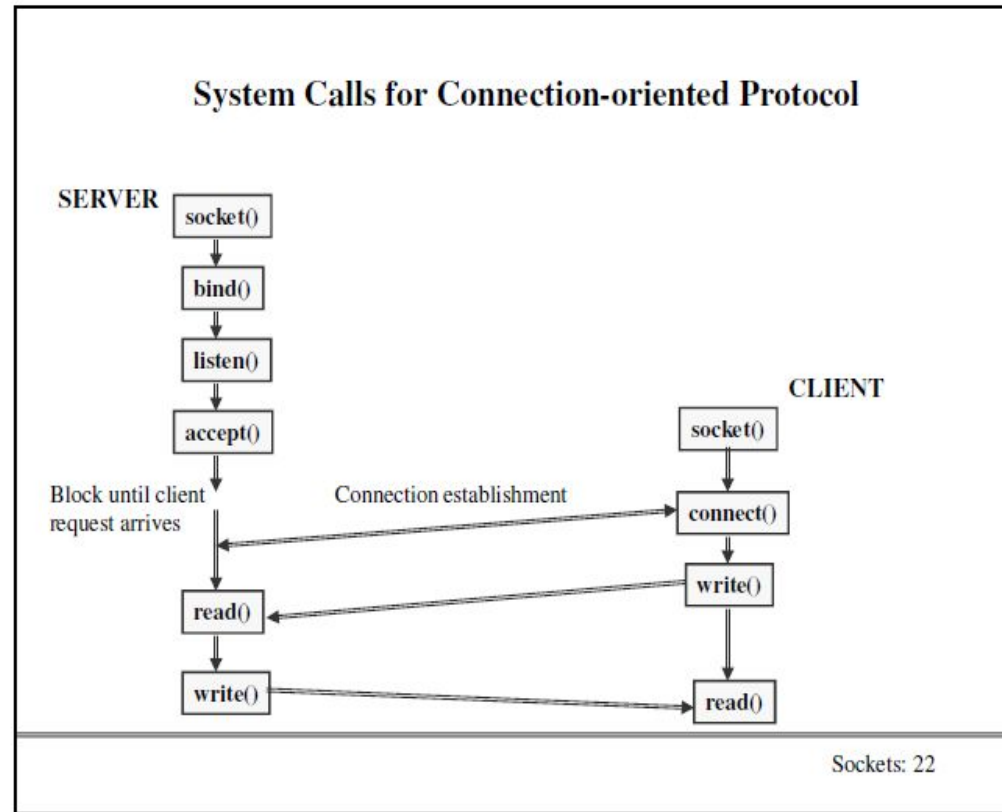
- When the server process has finished providing its service to the client, the server goes back to sleep, waiting for the next client request to arrive.

# Client-server Model (contd.)

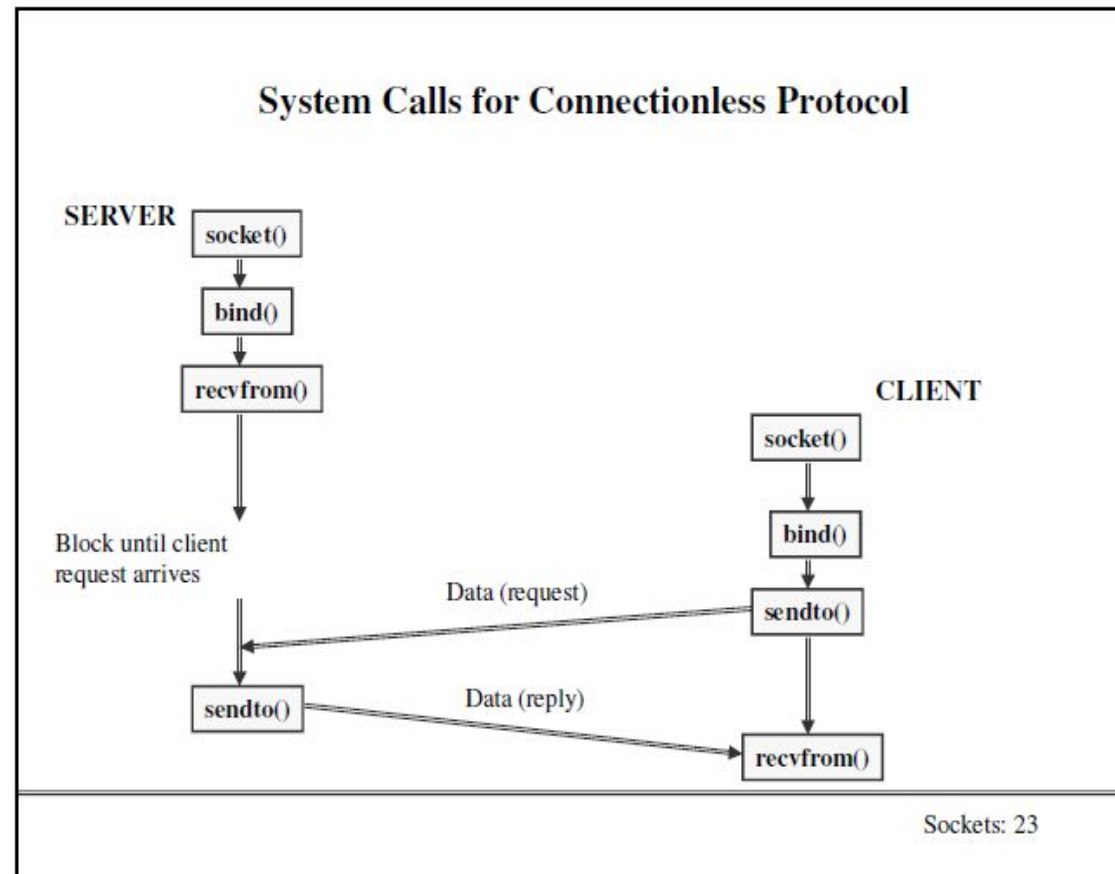
Roles of the client and the server processes are asymmetric.

- Two types of servers:
  - ***Iterative servers***: Used when the server process knows in advance how long it takes to handle each request and it handles each request itself.
  - ***Concurrent servers***: Used when the amount of work required to handle a request is unknown; the server starts another process to handle each request.

# System calls for Connection-oriented Protocol



# System calls for Connectionless Protocol



# References

- **Unix Network Programming**

*W.R.Stevens*, Prentice-Hall of India, 1992.

- **Internetworking with TCP/IP (Volume I,II,III)**

*D.E.Comer and D.L.Stevens*, Prentice-Hall of India, 1995.

- **<http://www.ecst.csuchico.edu/~beej/guide/net>**